

# ASSIGNMENT - 1

Group Number : C19

Group members : Bokka Srividya - 200101026

Yenugonda Suma Sree - 200101114

Jeram Niharika - 200101046

## Part 1: Kernel Threads

**Aim :** To implement kernel threads by adding system calls , specifically `thread_create` , `thread_join` , `thread_exit` without any process synchronization methods.

To create the system calls required , following files were edited :

**syscall.h :** Here we assign a unique number to every system call in the xv6 system. Before adding the new syscalls there were already 22 system calls. Now we add

- `#define SYS_thread_create 23`
- `#define SYS_thread_join 24`
- `#define SYS_thread_exit 25`

**syscall.c :** It stores an array of pointers to functions (`syscalls[]`) which use indexes defined in `syscall.h` to point to a respective system call function stored at a different memory location.

- `extern int sys_thread_create(void);`
- `extern int sys_thread_join(void);`
- `extern int sys_thread_exit(void);`

We also put a function prototype here (but not implementation).

- `[SYS_thread_create] sys_thread_create,`
- `[SYS_thread_create] sys_thread_join,`
- `[SYS_thread_create] sys_thread_exit.`

**sysproc.c:** Here the implementations of system calls are written.

**user.h** and **usys.S:** These are the interfaces for the system to access the system calls. The function prototype is added in `user.h` (included as header file in our program) while instruction to treat it as a system call is included in `usys.S`

**user.h :**

- `int thread_create(void()(void), void*, void*);`
- `int thread_join(void);`
- `void thread_exit(void);`

**usys.S :**

- `SYSCALL(thread_create)`
- `SYSCALL(thread_join)`
- `SYSCALL(thread_exit)`

In **sysproc.c** the implementation of system calls we just called the function .

**thread\_create**, **thread\_join** and **thread\_exit** are implemented in **proc.c** and declared in **defs.h**.

**thread\_create():**

We implemented `thread_create` similar to **fork()**, but in `fork()` the allocation of address space to the child process is different from the parent process and in our `thread_create` we gave the **same address** for the new process.

**myproc()** returns the currently running process in the cpu.

*struct proc \*curproc = myproc();*

Initially we create a new process we give it a *pid* for it and change the state of the process to **EMBRYO**.

*np->pgdir = curproc->pgdir;*

*np->tf->eip = (uint)fcn;*

We then set the instruction pointer to starting point of the function address

*np->tf->esp = (uint)stack+4096;*

And then we go to the end of the stack

Stack pointer go back by 4 bytes and place the arg at that address.

Other code in the `thread_create` is same as `fork` like duplicating the open file to the thread, current working directory and name of the process

After that it is changing the process state **EMBRYO** to **RUNNABLE**

**thread\_join():**

`thread_join` system call is used to make a process to wait while the other is in execution in cpu till its completion

`thread_join` returns the exited thread process id.

Initially it goes through the process table and searches for the children of the main process (which are threads here) and checks them if they are exited or not . If we have children and any of them are not exited then it makes the main process to wait(sleep) until one thread terminates.

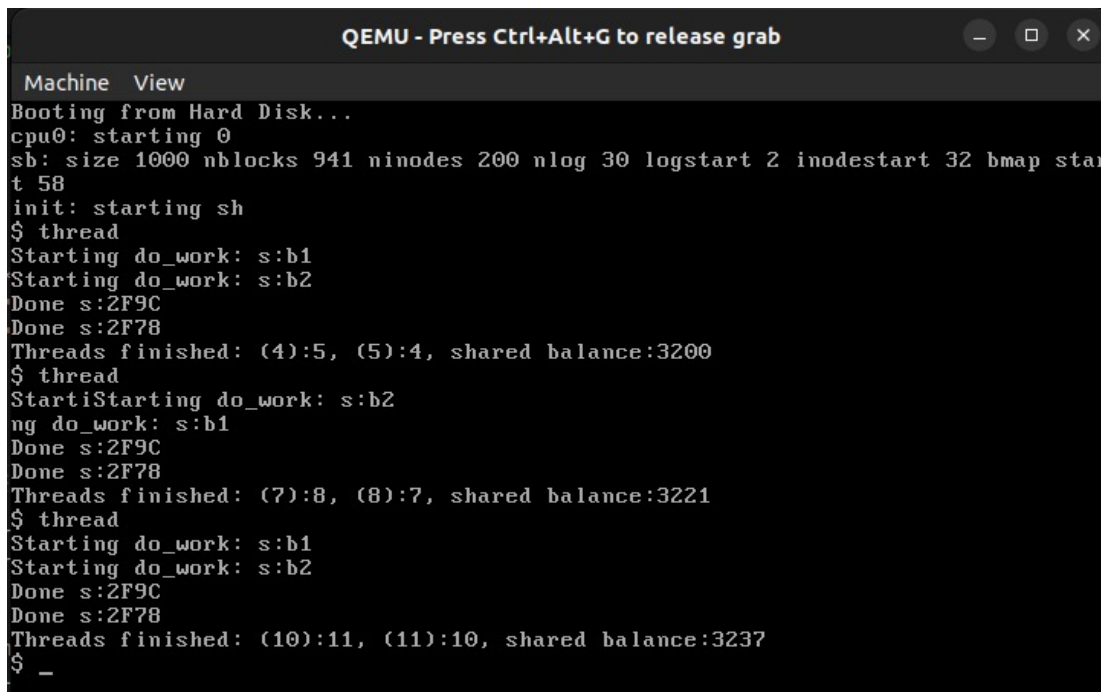
And we are not freeing the page directory (because main process also sharing same address space)

If thread is terminated then it reinitializes the process to an unused process and returns the pid.

### **thread\_exit():**

Fetches the currently running thread and close it all files and wakeup the parent process. If it has children processes then it will change the parent of that childrens to initproc. If any of that children is already terminated then it will wakeup the initproc set the state of the process to the ZOMBIE

**Observations :** Since there is no synchronization between the threads they may run into race conditions which are caused due to the shared variable (here the shared variable is “shared balance”) simultaneously .



```
QEMU - Press Ctrl+Alt+G to release grab
Machine View
Booting from Hard Disk...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap star
t 58
init: starting sh
$ thread
Starting do_work: s:b1
Starting do_work: s:b2
Done s:2F9C
Done s:2F78
Threads finished: (4):5, (5):4, shared balance:3200
$ thread
Starting do_work: s:b2
Starting do_work: s:b1
Done s:2F9C
Done s:2F78
Threads finished: (7):8, (8):7, shared balance:3221
$ thread
Starting do_work: s:b1
Starting do_work: s:b2
Done s:2F9C
Done s:2F78
Threads finished: (10):11, (11):10, shared balance:3237
$ _
```

- i) Here when initially threads were generated the balance of thread1 was 3200 and that of thread2 was 2800
- ii) The total shared balance to be shown was 6000 .
- iii) But we observe that the shared balance shown is less than 6000
- iv) This is because the threads run into race condition due to the shared variable – “shared balance”.

## Part 2 : Synchronization

**Aim :** To fix the synchronization error using spinlock and mutex primitives so that the threads would execute atomically

### Description :

#### 1)Spinlock:

We implemented spinlocks in the file thread.c in which we defined a thread\_spinlock data structure that has the attribute “**locked**”, which indicates when the lock is acquired or not.

**i) void thread\_spin\_init(struct thread\_spinlock \*lk):** This function initializes the lock to the correct initial state. We initialized locked attribute in thread\_spinlock to zero. When “locked” is equal to zero states that the thread\_spinlock is not acquired.

**ii) void thread\_spin\_lock(struct thread\_spinlock \*lk):**

In this, our implementation is such that we acquire the lock after using an atomic instruction xchg(Exchange). If the returned value of this xchg function is not equal to zero we enter into busy waiting. Else we use the atomic instruction “\_\_sync\_synchronize()” that tells the C compiler and the processor to not move loads or stores past this point, to ensure that the critical section’s memory references happen after the lock is acquired.

**iii) void thread\_spin\_unlock(struct thread\_spinlock \*lk):**

In this we use “\_\_sync\_synchronize()” and then we release the lock, equivalent to setting locked=0. We used asm volatile to set this locked value to zero because this code can’t use a C assignment as it might not be atomic.

#### 2)Mutexes:

Mutexes can be implemented similar to spinlocks except that it keeps the other processes or threads which are not in execution to sleep rather than busily waiting.

**i) void thread\_mutex\_init(struct thread\_mutexlock \*lk):**

This function initializes the lock to the correct initial state. We initialized locked attribute in thread\_mutexlock to zero.

When “locked” is equal to zero states that the thread\_mutexlock is not acquired.

**ii) void thread\_mutex\_lock(struct thread\_mutexlock \*lk):**

In this, our implementation is such that we acquire the lock after using an atomic instruction xchg(Exchange). If the returned value of this xchg function is not equal to zero that is when the lock is already acquired, the thread can sleep while waiting for the lock to be released. Else we use the atomic instruction “\_\_sync\_synchronize()” that tells the C compiler and the processor to not move loads or stores past this point, to ensure that the critical section’s memory references happen after the lock is acquired.

**iii) void thread\_mutex\_unlock(struct thread\_mutexlock \*lk):**

In this we use “\_\_sync\_synchronize()” and then we release the lock, equivalent to setting locked=0. We used asm volatile to set this locked value to zero because this code can’t use a C assignment as it might not be atomic.

**Observations :** Since we used the synchronization primitives the shared balance is used by only one thread which locks the shared variable, until the thread is executed and releases the lock the other thread cannot access the shared variable.

```
QEMU
Machine View
Booting from Hard Disk...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ thread
Starting do_work: s:b1
Starting do_work: s:b2
Done s:2F9C
Done s:2F78
Threads finished: (4):5, (5):4, shared balance:6000
$ thread
Starting do_work: s:b1
Starting do_work: s:b2
Done s:2F78
Done s:2F9C
Threads finished: (7):7, (8):8, shared balance:6000
$ thread
Starting do_work: s:b1
Starting do_work: s:b2
Done s:2F9C
Done s:2F78
Threads finished: (10):10, (11):11, shared balance:6000
$ _
```

Fig 01 : Threads execution with spinlocks

```
QEMU
Machine View
Booting from Hard Disk...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ thread
Starting do_work: s:b2
Starting do_work: s:b1
Done s:2F78
Done s:2F9C
Threads finished: (4):4, (5):5, shared balance:6000
$ thread
Starting do_work: s:b2
Starting do_work: s:b1
Done s:2F9C
Done s:2F78
Threads finished: (7):8, (8):7, shared balance:6000
$ thread
Starting do_work: s:b1
Starting do_work: s:b2
Done s:2F78
Done s:2F9C
Threads finished: (10):10, (11):11, shared balance:6000
$ _
```

Fig 02 : Threads execution with mutexes